



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

Comparison of GPU programming models for computational fluid dynamics

Robert Watt , Peter Jacobs , Rowan Gollan , and Shahzeb Imran
The University of Queensland

Australasian Fluid Mechanics Conference
December 02, 2024

The fragmentation of GPUs

Gadi: Nvidia




Figure 1: <https://www.nvidia.com/en-au/data-center/h100/>

Setonix: AMD Mi250X




Figure 2: <https://www.nvidia.com/en-au/data-center/h100/>

The fragmentation of GPUs

Gadi: Nvidia




Figure 6: <https://www.nvidia.com/en-au/data-center/h100/>

Setonix: AMD Mi250X




Figure 7: <https://www.nvidia.com/en-au/data-center/h100/>

Frontier: AMD Mi250X




Figure 8: <https://www.amd.com/en/products/accelerators/instinct/mi200/mi250x.html>

El Capitan: AMD Mi300




Figure 9: <https://www.amd.com/en/products/accelerators/instinct/mi300.html>

Aurora: Intel Data Center GPU Max



Figure 10: <https://www.tweaktown.com/news/89438/intels-new-data-center-gpu-max-1100-uses-controversial-12vhpwr-connector/index.html>

GPU programming models

CUDA

- Nvidia's language, for programming Nvidia GPUs
- Based on C/C++ or Fortran
- Runs on Nvidia GPUs

ROCm/HIP

- AMD's language for programming AMD GPUs
- Based on C/C++ or Fortran
- Runs on AMD and Nvidia GPUs

OpenMP

- Parallelism library built into C++ compilers
- Allows for offload to many devices, including GPUs

SYCL

- Provides an abstraction for heterogeneous programming in C++
- Implemented by various compilers (e.g. OneAPI by Intel)
- Runs on any (supported) accelerator

Kokkos

- C++ library providing abstractions for parallel programming
- Originally developed by Sandia National Lab, now a Linux Foundation project
- Runs on any supported accelerator
- Used in Trillinos

OpenCL

- Maybe dead? Maybe not?

RAJA

- Similar to Kokkos, but split into more libraries
- Runs on any (supported) accelerator

OCCA

- Used in the first trillion degree of freedom simulation
- Runs on any (supported) accelerator

Chapel

- Programming language from Cray/HP for high-performance computing
- Runs on CPUs, Nvidia GPUs, with some AMD and Intel GPU support

julia

- Programming language with libraries to execute on Nvidia and AMD GPUs

That's a lot of options...

Which programming model should we use?

That's a lot of options...

Which programming model should we use?

How well do they run across the different GPU brands?

That's a lot of options...

Which programming model should we use?

How well do they run across the different GPU brands?

Today's talk: Compare Kokkos with pure CUDA

Kokkos overview

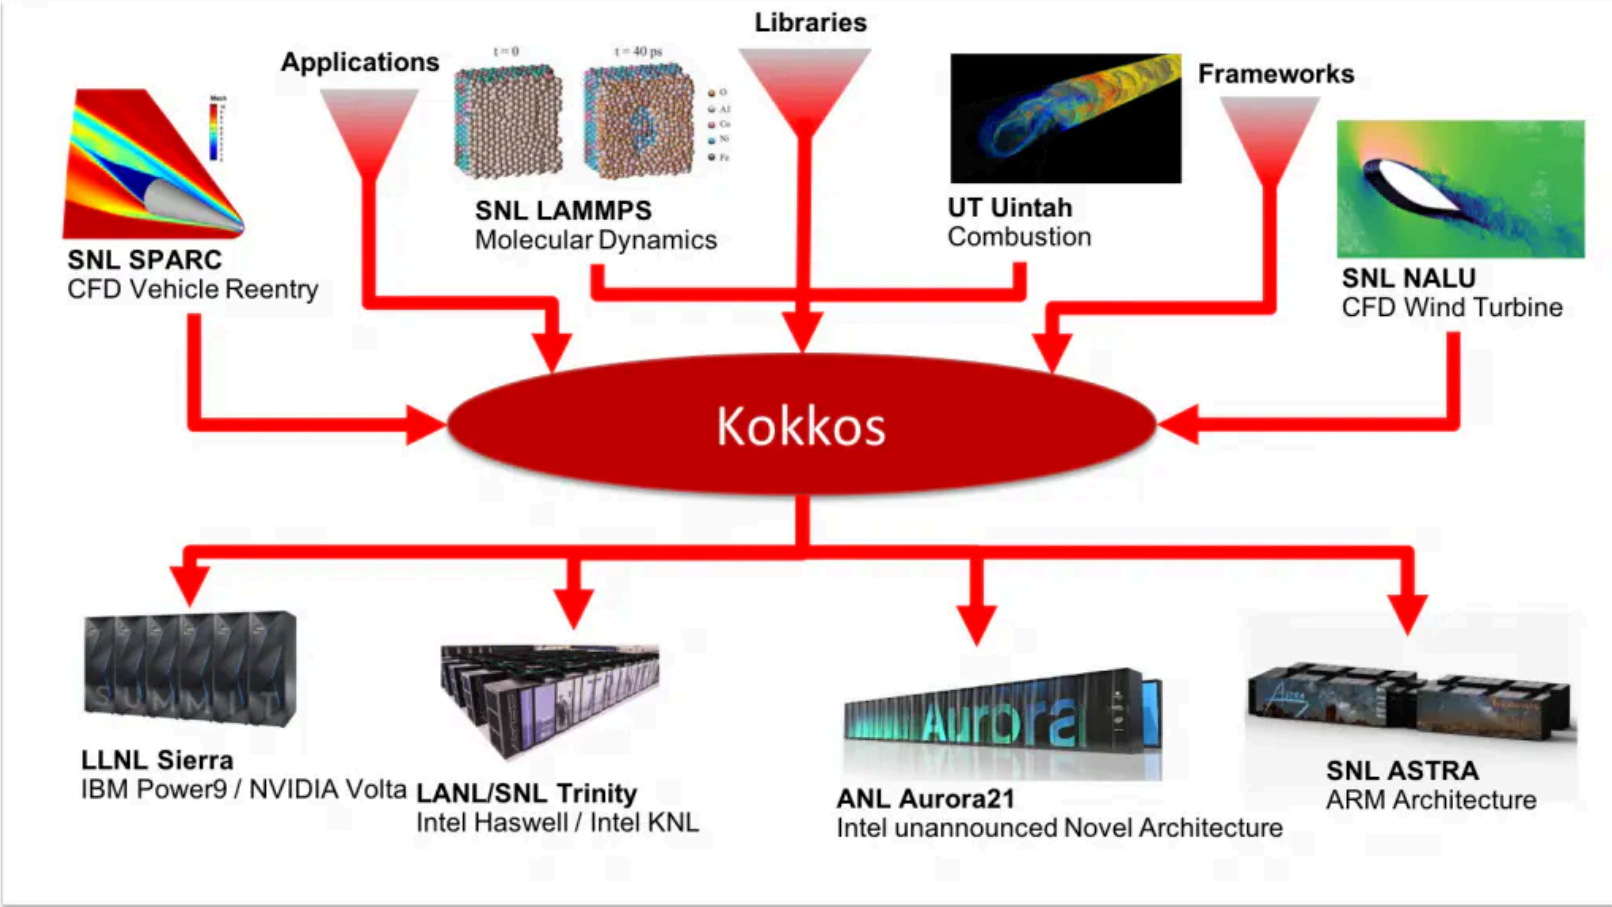


Figure 11: Kokkos architecture (<https://kokkos.org/about/abstract/>)

Kokkos overview

- C++ library managing mapping of parallelism of the algorithm to the parallelism of the hardware
- Can compile using vendor-provided compilers (e.g. CUDA or HIP), or can use SYCL, OpenMP, C++ threads etc.
- Kokkos manages shared memory parallelisation
 - Compatible with MPI, HPX, PGAS etc.

Advantages

- Automatic memory management
- Compile-time configurable memory layout (row major vs column major arrays)
 - Helps get the best memory layout for a particular piece of hardware
- Transparent management of different memory spaces
- Can perform device specific optimisations
- Free and open-source

Disadvantages

- Still C++
- Relies heavily on template meta-programming
 - may be confusing for C++ beginners
 - Compile times are longer, but not necessarily prohibitive

Example program

CUDA

```
#ifdef HIP
#define cudaMalloc hipMalloc
#define cudaFree hipFree
#endif

__global__
void initialise(double* a, int size) {
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < size) {
        a[i] = 0.0;
    }
}

__global__
void add_one(double* a, int size) {
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < size) {
        a[i] += 1.0;
    }
}

int main() {
    // allocate an array of 10 doubles on the GPU
    double* a;
    cudaMalloc(&a, 10*sizeof(double));
    initialise<<<1, 32>>>>(a, 10);
    cudaDeviceSynchronize();
    add_one<<<1, 32>>>>(a, 10);
    cudaDeviceSynchronize();
    cudaFree(a);
}
```

Kokkos

```
#include <Kokkos_Core.hpp>
int main() {
    // allocate an array of 10 doubles
    // on the default device
    // (these will be zero initialised)
    Kokkos::View<double*, KokkosDefaultMemSpace> a(10);

    // Add one to each double in parallel
    Kokkos::parallel_for(10, KOKKOS_LAMBDA(const int i){
        a(i) += 1.0;
    });
}
```

Chicken 🐔

- CUDA/HIP
- Block-structured grids
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Van-Albada limiter
- Viscous gradients calculated with least-squares
- Time integration:
 - 3rd order Runge-Kutta time integration

The codes

Chicken 🐔

- CUDA/HIP
- Block-structured grids
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Van-Albada limiter
- Viscous gradients calculated with least-squares
- Time integration:
 - 3rd order Runge-Kutta time integration

```
MM      --
<' \___/|  --
  \___/  --
    ] [  --
-----
CHICKEN
-----
```

The codes

Chicken 🐔

- CUDA/HIP
- Block-structured grids
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Van-Albada limiter
- Viscous gradients calculated with least-squares
- Time integration:
 - 3rd order Runge-Kutta time integration

```
MM      --
<' \___/|  --
  \___/  --
    ] [  --
-----
CHICKEN
-----
```

Ibis 🗑️🐔

- Kokkos
- Unstructured grids
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Barth-Jespersen limiter
- Inviscid and viscous gradients calculated with least-squares
- Time integration:
 - Any explicit Runge-Kutta
 - Jacobian-Free Newton-Krylov

The codes

Chicken 🐔

- CUDA/HIP
- **Block-structured grids**
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Van-Albada limiter
- Viscous gradients calculated with least-squares
- Time integration:
 - 3rd order Runge-Kutta time integration

```
MM      --
<' \___/|  --
  \___/  --
    ] [  --
-----
CHICKEN
-----
```


Ibis 🗑️ 🐔

- Kokkos
- **Unstructured grids**
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Barth-Jespersen limiter
- **Inviscid** and viscous **gradients** calculated with least-squares
- Time integration:
 - Any explicit Runge-Kutta
 - Jacobian-Free Newton-Krylov

The codes


Chicken 🐔

- CUDA/HIP
- **Block-structured grids**
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Van-Albada limiter
- Viscous gradients calculated with least-squares
- Time integration:
 - 3rd order Runge-Kutta time integration



Ibis 🗑️ 🐔

- Kokkos
- **Unstructured grids**
- Finite volume
 - Selection of upwind flux calculators
 - 2nd order accuracy via MUSCL style reconstruction
 - Barth-Jespersen limiter
- **Inviscid** and viscous **gradients** calculated with least-squares
- Time integration:
 - Any explicit Runge-Kutta
 - Jacobian-Free Newton-Krylov



The test problem

The problem

- Gaseous injection into Mach 4 cross flow




Figure 12: Domain and boundary conditions

Conditions

	Pressure (kPa)	Temperature (K)	Velocity (m/s)
Free-stream	1.013	300	1390
Injector	10.013	300	350

Grid

- Various grid resolutions from 0.5-3.5 million cells
- Chicken used blocking structure indicated in Figure 12
- Ibis merged all blocks into one block

Numerics

- AUSMDV flux calculator
- 2nd order accurate spatially
- 3rd order Runge-Kutta time-stepping

The test problem




Figure 13: Flow field after 5ms

Code performance with grid resolution

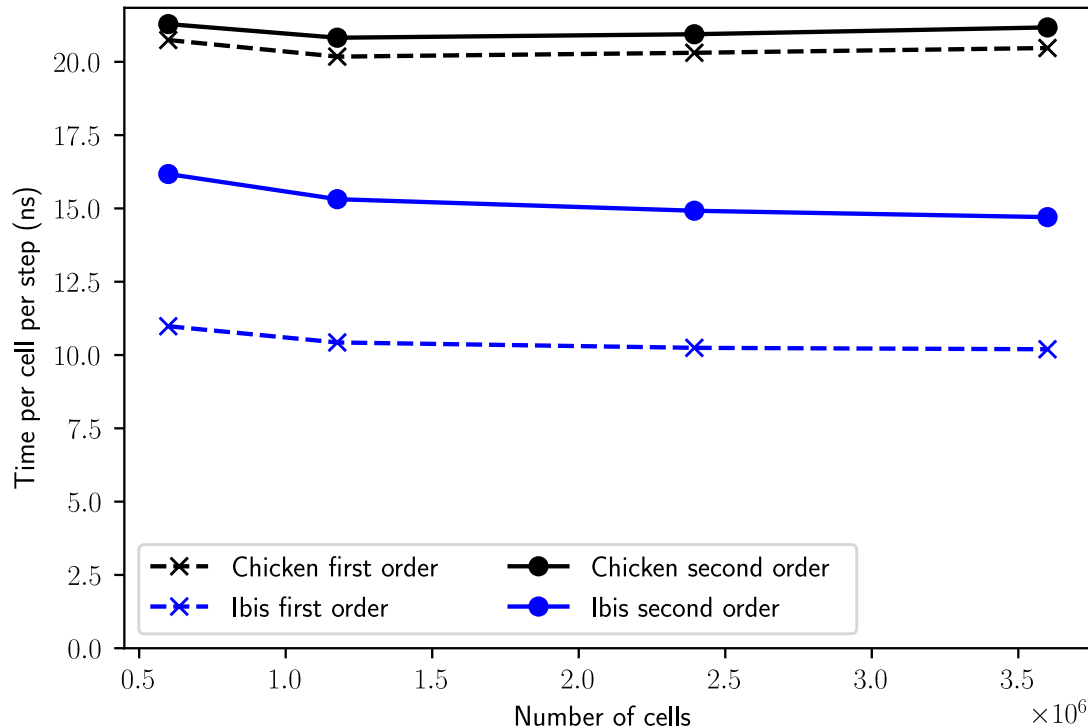


Figure 14: Time to update a cell at various grid resolutions on an Nvidia H100

Motivation

- Check to make sure the GPU is saturated
 - Chicken may require more cells to saturate the GPU
- Get base-line performance of each code

Reasons for differences

- Chicken moves memory for 2nd order around even when doing 1st order, so 1st order is slow
- Chicken calculates all fluxes in one large kernel
 - This may be causing register pressure, leading to lower occupancy than Ibis
 - The kernel with the lowest occupancy in Ibis had twice the occupancy of Chicken's fused kernel
- Chicken uses array-of-structures, Ibis uses structure-of-arrays

Acceleration on different hardware

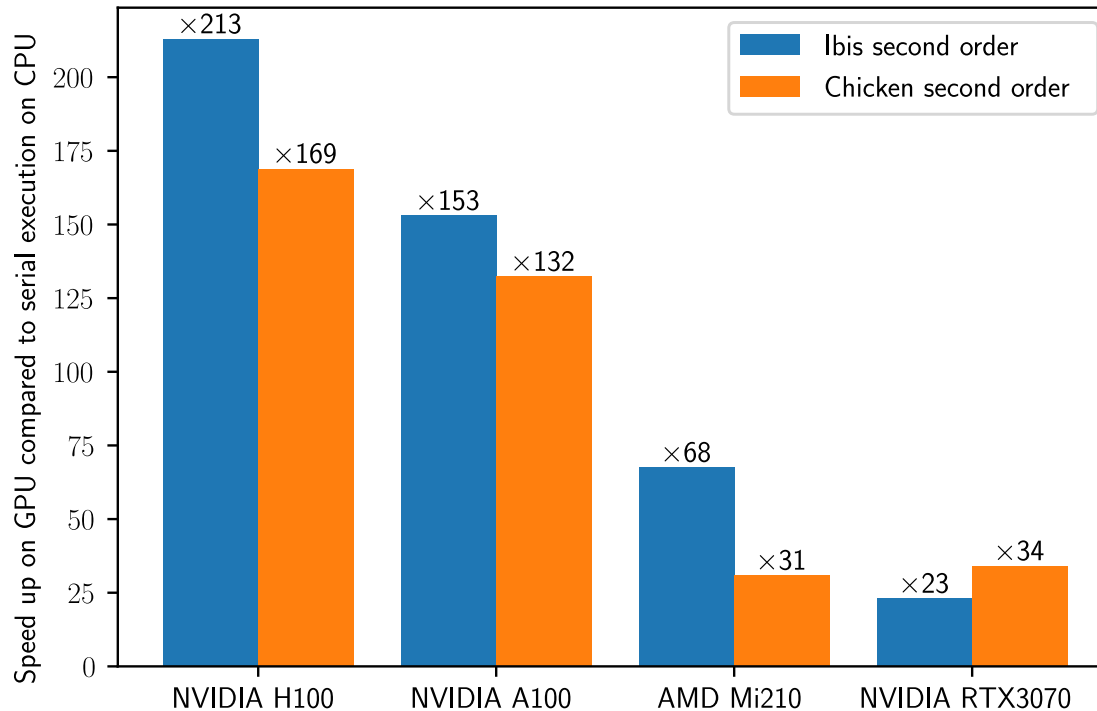



Figure 15: Acceleration compared to serial CPU performance on various GPUs

- Ibis (Kokkos) had greater acceleration on all HPC grade GPUs
 - Likely due to Ibis having greater occupancy owing to smaller kernels, smaller register count, higher occupancy, and better memory access patterns
- Chicken (CUDA) had greater acceleration on consumer grade GPU
 - Likely due to few FP64 cores on consumer grade GPUs

GPU utilisation



- Closer to the black bline is better (above is impossible)
- Some long-running kernels in Ibis benefit from being their own kernel with lower register pressure
- Ibis does more floating point operations per byte of data transferred partially due to better memory layout.

Figure 16: Roofline plot of the two codes on an Nvidia A100. Size of the circles represents time spent in that kernel

Conclusions

- Writing a code in CUDA doesn't mean its fast
 - Leave the details of the parallelism to the experts, and we'll write the physics
- Code design and memory access patterns had a larger impact on performance than overhead from abstractions
- Kokkos seems to be a good option for performance portable code
- But maybe one of the others is better?

```
MM
<' \__ /|
  \_ _/
    ] [
  __ /- \__
|-----|
| | | | |
| | | | |
| | | | |
| | | | |
|-----|
```

```
MM      --
<' \__ /|  --
  \_ _/    --
    ] [    --
-----
CHICKEN
-----
```